

Improving Area and Resoudesign.rce Utilization Lab

Introduction

This lab introduces various techniques and directives which can be used in Vitis HLS to improve design performance as well as area and resource utilization. The design under consideration performs Discrete Cosine Transformation (DCT) on an 8x8 block of data.

This design implements a discrete cosine transformation (DCT), and it is provided as C source. The function leverages a 2D DCT algorithm by first processing each row of the input array via a 1D DCT, then processing the columns of the resulting array through the same 1D DCT. It calls the `read_data`, `dct_2d`, and `write_data` functions.

Objectives

After completing this lab, you will be able to:

- Add directives in your design
- Improve performance using PIPELINE directive
- Distinguish between DATAFLOW directive and Configuration Command functionality
- Apply memory partitions techniques to improve resource utilization

Steps

Validate the Design from Command Line

Validate your design in the terminal.

1. Open a terminal.
2. Change directory to `{labs}/lab3`.

A self-checking program (`dct_test.c`) is provided. Using that we can validate the design. A Makefile is also provided. Using the Makefile, the necessary source files can be compiled and the compiled program can be executed. 3. Invoke Vitis HLS Command prompt by selecting **Start > Xilinx Design Tools > Vitis HLS 2021.2 Command Prompt**. In the terminal, type `make` to compile and execute the program.

```
gcc -ggdb -w -I/include -c -o dct.o dct.c
gcc -lm -lstdc++ dct.o dct_test.o -o dct
./dct
*** *** ***
Results are good
*** *** *** ***
```

Validating the design

Note that the source files (`dct.c` and `dct_test.c` are compiled, then `dct` executable program was created, and then it was executed. The program tests the design and outputs `Results are good` message.

Create a New Project

Create a new project in Vitis HLS GUI targeting xc7z020clg400-1.

1. Type **vitis_hls** in the terminal.
2. In the Vitis HLS GUI, click on **Create Project**. The **New Vitis HLS Project** wizard opens.
3. Click **Browse...** button of the Location field and browse to **{labs}/lab3** and then click **OK**.
4. For Project Name, type **dct** and click **Next**.
5. In the Add/Remove Files for the source files, type **dct** as the top function name (the provided source file contains the function, to be synthesized, called **dct**).
6. Click the **Add Files...** button, select **dct.c** file from the **{labs}/lab3** folder, and then click **Open**.
7. Click **Next**.
8. In the *Add/Remove Testbench Files* for the testbench, click the **Add Files...** button, select **dct_test.c**, **in.dat**, **out.golden.dat** files from the **{labs}/lab3** folder and click **Open**.
9. Click **Next**.
10. In the *Solution Configuration* page, leave *Solution Name* field as **solution1** and set the clock period as **10**. Leave *Uncertainty* field blank.
11. Click on *Part's* Browse button, and select the following filters, using the Parts Specify option, to select **xc7z020clg400-1**.
12. Click **Finish**.
13. Double-click on the **dct.c** under the source folder to open its content in the information pane.


```

104 void dct(short input[N], short output[N])
105 {
106
107     short buf_2d_in[DCT_SIZE][DCT_SIZE];
108     short buf_2d_out[DCT_SIZE][DCT_SIZE];
109
110     // Read input data. Fill the internal buffer.
111     read_data(input, buf_2d_in);
112
113     dct_2d(buf_2d_in, buf_2d_out);
114
115     // Write out the results.
116     write_data(buf_2d_out, output);
117 }
118

```

The design under consideration

The top-level function *dct*, is defined at line 104. It implements 2D DCT algorithm by first processing each row of the input array via a 1D DCT then processing the columns of the resulting array through the same 1D DCT. It calls `read_data`, `dct_2d`, and `write_data` functions. The `read_data` function is defined at line 80 and consists of two loops – RD_Loop_Row and RD_Loop_Col. The `write_data` function is defined at line 92 and consists of two loops to perform writing the result. The `dct_2d` function, defined at line 49, calls `dct_1d` function and performs transpose. Finally, `dct_1d` function, defined at line 30, uses `dct_coeff_table` and performs the required function by implementing a basic iterative form of the 1D Type-II DCT algorithm. Following figure shows the function hierarchy on the left-hand side, the loops in the order they are executes and the flow of data on the right-hand side.

*Design hierarchy and dataflow*

Synthesize the Design

Synthesize the design with the defaults. View the synthesis results and answer the question listed in the detailed section of this step.

1. Select **Solution > Run C Synthesis > Active Solution** to start the synthesis process.
2. When synthesis is completed, several report files will become accessible and the *Synthesis Results* will be displayed in the information pane.

The `dct_1d`, `dct_2d`, `read_data` and `write_data` functions are inlined. Verify this by scrolling up into the Vitis HLS Console view.

```
INFO: [HLS 214-178] Inlining function 'dct_1d' into 'dct_2d' (dct.c:25:0)
INFO: [HLS 214-178] Inlining function 'read_data' into 'dct' (dct.c:79:0)
INFO: [HLS 214-178] Inlining function 'dct_2d' into 'dct' (dct.c:79:0)
INFO: [HLS 214-178] Inlining function 'write_data' into 'dct' (dct.c:79:0)
```

Inlining of dct_1d, dct_2d, read_data and write_data functions

3. The *Synthesis Report* shows the performance and resource estimates as well as estimated latency in the design. Note that the design is already pipelined.

The screenshot displays the Synthesis Report interface with three main sections:

- Performance:** Shows clock summary for ap_clk: Target 10.00 ns, Estimated 6.508 ns, Uncertainty 2.70 ns.
- Latency:** Shows latency summary for a single instance: min 423 cycles, max 423 cycles, absolute min 4.230 us, absolute max 4.230 us, interval min 424 cycles, interval max 424 cycles, type no.
- Utilization Estimates:** Shows resource summary for various components:

Name	BRAM_18K	DSP	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	-	-	-
FIFO	-	-	-	-	-
Instance	-	16	792	1396	-
Memory	17	-	119	16	0
Multiplexer	-	-	-	600	-
Register	-	-	18	-	-
Total	17	16	929	2012	0
Available	280	220	106400	53200	0
Utilization (%)	6	7	~0	3	0

Synthesis report

4. Using scroll bar on the right, scroll down into the report and answer the following question.

Question 1

Answer the following question:

Estimated clock period:

Worst case latency:

Number of DSP48E used:

Number of BRAMs used:

Number of FFs used:

Number of LUTs used:

5. The report also shows the top-level interface signals generated by the tools.

Interface						
Summary						
RTL Ports	Dir	Bits	Protocol	Source Object	C Type	
ap_local_block	out	1	ap_ctrl_hs	dct	return value	
ap_local_deadlock	out	1	ap_ctrl_hs	dct	return value	
ap_clk	in	1	ap_ctrl_hs	dct	return value	
ap_rst	in	1	ap_ctrl_hs	dct	return value	
ap_start	in	1	ap_ctrl_hs	dct	return value	
ap_done	out	1	ap_ctrl_hs	dct	return value	
ap_idle	out	1	ap_ctrl_hs	dct	return value	
ap_ready	out	1	ap_ctrl_hs	dct	return value	
input_r_address0	out	6	ap_memory	input_r	array	
input_r_ce0	out	1	ap_memory	input_r	array	
input_r_q0	in	16	ap_memory	input_r	array	
output_r_address0	out	6	ap_memory	output_r	array	
output_r_ce0	out	1	ap_memory	output_r	array	
output_r_we0	out	1	ap_memory	output_r	array	
output_r_d0	out	16	ap_memory	output_r	array	

Generated interface signals

You can see ap_clk, ap_rst are automatically added. The ap_start, ap_done, ap_idle, and ap_ready are top-level signals used as handshaking signals to indicate when the design is able to accept next computation command (ap_idle), when the next computation is started (ap_start), and when the computation is completed (ap_done). The top-level function has input and output arrays, hence an ap_memory interface is generated for each of them.

Run Co-Simulation

Run the Co-simulation, selecting Verilog. Verify that the simulation passes.

1. Select **Solution > Run C/RTL Co-simulation** to open the dialog box so the desired simulations can be run. A C/RTL Co-simulation Dialog box will open.
2. Select the **Verilog** option, and click **OK** to run the Verilog simulation using Vivado XSIM simulator.


The RTL Co-simulation will run, generating and compiling several files, and then simulating the design. In the console window you can see the progress and also a message that the test is passed.

```
INFO: [COSIM 212-316] Starting C post checking ...
*** *** ***
Results are good
*** *** ***
INFO: [COSIM 212-1000] *** C/RTL co-simulation finished: PASS ***
```

RTL Co-Simulation results


Remove the pipeline optimization done by Vitis HLS automatically by adding pipeline off pragma

1. Select **Project > New Solution**.
2. A *Solution Configuration* dialog box will appear. Note that the check boxes of *Copy directives and constraints from solution* are checked with *solution1* selected. Click the **Finish** button to create a new solution with the default settings.




Creating a new Solution after copying the existing solution

3. Make sure that the **dct.c** source is opened and visible in the information pane, and click on the **Directive** tab.
4. Select function **DCT_Inner_Loop** in the directives pane, right-click on it, and select **Insert Directive...**
5. Click on the drop-down button of the *Directive* field. A pop-up menu shows up listing various directives. Select **PIPELINE** directive.
6. In the *Vitis HLS Directive Editor* dialog box, click on the **off** option to turn off the automatic pipelining. Make sure that the *Directive File* is selected as destination. Click **OK**.



Add PIPELINE off directive

7. Similarly, apply the **PIPELINE off** directive to **DCT_Outer_Loop**, **Row_DCT_Loop**, **Xpose_Row_Outer_Loop**, **Xpose_Row_Inner_Loop**, **Col_DCT_Loop**, **Xpose_Col_Outer_Loop**, **Xpose_Col_Inner_Loop**, **RD_Loop_Row**, **RD_Loop_Col**, **WR_Loop_Row**, and **WR_Loop_Col** objects. At this point, the *Directive* tab should look like as follows.



PIPELINE off directive applied

8. Click on the **Synthesis** button.
9. When the synthesis is completed, report shows the performance and area without the automatic optimization of Vitis HLS.


Timing Estimate															
Target	Estimated	Uncertainty													
10.00 ns	6.508 ns	2.70 ns													
Performance & Resource Estimates															
<input type="checkbox"/> Modules	<input checked="" type="checkbox"/> Loops	<input type="checkbox"/>													
Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ dct					5990	5.990E4	-	5991	-	no	5	2	286	916	0
▼ RD_Loop_Row					144	1.440E3	18	-	8	no	-	-	-	-	-
RD_Loop_Col					16	160.000	2	-	8	no	-	-	-	-	-
▼ Row_DCT_Loop					2704	2.704E4	338	-	8	no	-	-	-	-	-
DCT_Outer_Loop					336	3.360E3	42	-	8	no	-	-	-	-	-
▼ Xpose_Row_Outer_Loop					144	1.440E3	18	-	8	no	-	-	-	-	-
Xpose_Row_Inner_Loop					16	160.000	2	-	8	no	-	-	-	-	-
▼ Col_DCT_Loop					2704	2.704E4	338	-	8	no	-	-	-	-	-
DCT_Outer_Loop					336	3.360E3	42	-	8	no	-	-	-	-	-
▼ Xpose_Col_Outer_Loop					144	1.440E3	18	-	8	no	-	-	-	-	-
Xpose_Col_Inner_Loop					16	160.000	2	-	8	no	-	-	-	-	-
▼ WR_Loop_Row					144	1.440E3	18	-	8	no	-	-	-	-	-
WR_Loop_Col					16	160.000	2	-	8	no	-	-	-	-	-

Performance after applying PIPELINE off directive

Apply PIPELINE Directive

Create a new solution by copying the previous solution settings. Apply the PIPELINE directive to DCT_Inner_Loop, Xpose_Row_Inner_Loop, Xpose_Col_Inner_Loop, RD_Loop_Col, and WR_Loop_Col. Generate the solution and analyze the output.

1. Select **Project > New Solution**.
2. A *Solution Configuration* dialog box will appear. Click the **Finish** button (with copy from Solution2 selected).
3. Make sure that the **dct.c** source is opened in the information pane and click on the **Directive** tab.
4. Select the pragma HLS PIPELINE off of **DCT_Inner_Loop** of the **dct_1d** function in the *Directive* pane, right-click on it and select **Modify Directive**
5. In the Vitis HLS Directive Editor dialog box, click the **off** option to turn on the pipelining. Make sure that the *Directive File* is selected as destination. Click OK.
6. Leave II (Initiation Interval) blank as Vitis HLS will try for an II=1, one new input every clock cycle.
7. Click **OK**.
8. Similarly, apply the **PIPELINE** directive to **Xpose_Row_Inner_Loop** and **Xpose_Col_Inner_Loop** of the **dct_2d** function, and **RD_Loop_Col** of the **read_data** function, and **WR_Loop_Col** of the **write_data** function. But remove the **PIPELINE** directive of **DCT_Outer_Loop**, **Row_DCT_Loop**, **Xpose_Row_Outer_Loop**, **Col_DCT_Loop**, **Xpose_Col_Outer_Loop**, **RD_Loop_Row** and **WR_Loop_Row**. At this point, the Directive tab should look like as follows.



PIPELINE directive applied

9. Click on the **Synthesis** button.
10. When the synthesis is completed, select **Project > Compare Reports...** to compare the two solutions.
11. Select *Solution2* and *Solution3* from the *Available Reports*, click on the **Add>>** button, and then click **OK**.
12. Observe that the latency reduced from 5990 to 2451 clock cycles.

Performance Estimates			
Timing			
Clock		solution2	solution3
ap_clk	Target	10.00 ns	10.00 ns
	Estimated	6.508 ns	6.508 ns
Latency			
		solution2	solution3
Latency (cycles)	min	5990	2451
	max	5990	2451
Latency (absolute)	min	59.900 us	24.510 us
	max	59.900 us	24.510 us
Interval (cycles)	min	5991	2452
	max	5991	2452

Performance comparison after pipelining

13. Scroll down in the comparison report to view the resources utilization. Observe that the FFs and/or LUTs utilization increased whereas BRAM and DSP48E remained same.

Utilization Estimates		
	solution2	solution3
BRAM_18K	5	5
DSP	2	2
FF	286	515
LUT	916	1493
URAM	0	0

Resources utilization after pipelining

Open the Schedule Viewer and determine where most of the clock cycles are spend, i.e. where the large latencies are.

1. Click on the *Solution > Open Schedule Viewer*.
2. Select the **dct** entry and observe the **RD_Loop_Row_RD_Loop_Col** and **WR_Loop_Row_WR_Loop_Col** entries in the *Performance & Resource Estimates*. These are two nested loops flattened and given the new names formed by appending inner loop name to the outer loop name. You can verify this by looking in the *Console* view message.

```
INFO: [XFORM 203-541] Flattening a loop nest 'RD Loop_Row' '(dct,c:56:8)' in function 'dct'.
WARNING: [HLS 200-960] Cannot flatten loop 'DCT_Outer_Loop' '(dct,c:16:17)' in function 'dct' the outer loop is not a perfect loop because there is nontrivial logic in the loop latch.
Resolution: For help on HLS 200-960 see www.xilinx.com/cgi-bin/docs/doc?v=2021.2;t=hls+guidance;d=200-960.html
INFO: [XFORM 203-541] Flattening a loop nest 'Row_DCT_Loop' '(dct,c:28:13)' in function 'dct'.
INFO: [XFORM 203-541] Flattening a loop nest 'Xpose_Row_Outer_Loop' '(dct,c:28:16)' in function 'dct'.
WARNING: [HLS 200-960] Cannot flatten loop 'DCT_Outer_Loop' '(dct,c:16:17)' in function 'dct' the outer loop is not a perfect loop because there is nontrivial logic in the loop latch.
Resolution: For help on HLS 200-960 see www.xilinx.com/cgi-bin/docs/doc?v=2021.2;t=hls+guidance;d=200-960.html
INFO: [XFORM 203-541] Flattening a loop nest 'Col_DCT_Loop' '(dct,c:28:13)' in function 'dct'.
INFO: [XFORM 203-541] Flattening a loop nest 'Xpose_Col_Outer_Loop' '(dct,c:28:16)' in function 'dct'.
INFO: [XFORM 203-541] Flattening a loop nest 'WR_Loop_Row' '(dct,c:68:8)' in function 'dct'.
```

The console view content indicating loops flattening

3. In the *Performance & Resource Estimates* pane, expand all items. Notice that the most of the latency occurs is in **Row_DCT_Loop_DCT_Outer_Loop** and **Col_DCT_Loop_DCT_Outer_Loop** function.

			2451	2.451E4	-	2452	-
▼ dct							
▼ dct_Pipeline_RD_Loop_Row_RD_Loop_Col			-	67	670.000	-	67
🕒 RD_Loop_Row_RD_Loop_Col			-	65	650.000	3	1
▼ dct_Pipeline_Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop			-	67	670.000	-	67
🕒 Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop			-	65	650.000	3	1
▼ dct_Pipeline_Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop			-	67	670.000	-	67
🕒 Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop			-	65	650.000	3	1
▼ dct_Pipeline_WR_Loop_Row_WR_Loop_Col			-	67	670.000	-	67
🕒 WR_Loop_Row_WR_Loop_Col			-	65	650.000	3	1
🕒 Row_DCT_Loop_DCT_Outer_Loop			-	1088	1.088E4	17	-
▼ dct_Pipeline_DCT_Inner_Loop			-	13	130.000	-	13
🕒 DCT_Inner_Loop			-	11	110.000	5	1
🕒 Col_DCT_Loop_DCT_Outer_Loop			-	1088	1.088E4	17	-
▼ dct_Pipeline_DCT_Inner_Loop1			-	13	130.000	-	13
🕒 DCT_Inner_Loop			-	11	110.000	5	1

The Performance & Resource Estimates of dct

4. In the **Schedule Viewer** pane, select the **Col_DCT_Loop_DCT_Outer_Loop** entry, right-click on the **col_outbuf_addr_write_In19** (write) block in the **Schedule Viewer**, and select **Goto Source**. Notice that line 45 is highlighted which is preventing the flattening of the DCT_Outer_Loop.

```

38 DCT_Outer_Loop:
39   for (k = 0; k < DCT_SIZE; k++) {
40     DCT_Inner_Loop:
41       for(n = 0, tmp = 0; n < DCT_SIZE; n++) {
42         int coeff = (int)dct_coeff_table[k][n];
43         tmp += src[n] * coeff;
44       }
45       dst[k] = DESCALE(tmp, CONST_BITS);
46     }
47   }


```

Understanding what is preventing DCT_Outer_Loop flattening

5. Switch to the **Synthesis** perspective.

Create a new solution by copying the previous solution settings. Apply fine-grain parallelism of performing multiply and add operations of the inner loop of **dct_1d using **PIPELINE** directive by moving the **PIPELINE** directive from inner loop to the outer loop of **dct_1d**. Generate the solution and analyze the output.**

1. Select **Project > New Solution**.
2. A **Solution Configuration** dialog box will appear. Click the **Finish** button (with **Solution3** selected).
3. Select **PIPELINE** directive of **DCT_Inner_Loop** of the **dct_1d** function in the Directive pane, right-click on it and select **Remove Directive**.
4. Select **DCT_Outer_Loop** of the **dct_1d** function in the Directive pane, right-click on it and select **Insert Directive...**
5. A pop-up menu shows up listing various directives. Select **PIPELINE** directive.
6. Click **OK**.



PIPELINE directive applied to DCT_Outer_Loop

By pipelining an outer loop, all inner loops will be unrolled automatically (if legal), so there is no need to explicitly apply an UNROLL directive to DCT_Inner_Loop. Simply move the pipeline to the outer loop: the nested loop will still be pipelined but the operations in the inner-loop body will operate concurrently.

7. Click on the **Synthesis** button.
8. When the synthesis is completed, select **Project > Compare Reports...** to compare the two solutions.
9. Select *Solution3* and *Solution4* from the **Available Reports**, click on the **Add>>** button, and then click **OK**.
10. Observe that the latency reduced from 2451 to 643 clock cycles.

Performance Estimates			
Timing			
Clock		solution3	solution4
ap_clk	Target	10.00 ns	10.00 ns
	Estimated	6.508 ns	6.508 ns
Latency			
		solution3	solution4
Latency (cycles)	min	2451	643
	max	2451	643
Latency (absolute)	min	24.510 us	6.430 us
	max	24.510 us	6.430 us
Interval (cycles)	min	2452	644
	max	2452	644

Performance comparison after pipelining

11. Scroll down in the comparison report to view the resources utilization. Observe that the utilization of DSP and FF increased. Since the DCT_Inner_Loop was unrolled, the parallel computation requires 8 DSP48E.

	solution3	solution4
BRAM_18K	5	5
DSP	2	8
FF	515	827
LUT	1493	1377
URAM	0	0

Resources utilization after pipelining

Perform design analysis and look at the dct performance view.

1. Switch to the **Performance & Resource Estimates** in Synthesis Summary.
2. Expand, if necessary, and notice that the DCT_Outer_Loop is now pipelined and there is no DCT_Inner_Loop entry.

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▀ dct	-	-	-	-	643	6.430E3	-	644	-	no	5	8	827	1377	0
▀ dct_Pipeline_RD_Loop_Row_RD_Loop_Col	-	-	-	-	67	670.000	-	67	-	no	0	0	32	156	0
RD_Loop_Row_RD_Loop_Col	-	-	-	-	65	650.000	3	1	64	yes	-	-	-	-	-
▀ dct_Pipeline_Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	-	-	-	67	670.000	-	67	-	no	0	0	38	170	0
Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	-	-	-	65	650.000	3	1	64	yes	-	-	-	-	-
▀ dct_Pipeline_Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	-	-	-	67	670.000	-	67	-	no	0	0	38	170	0
Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	-	-	-	65	650.000	3	1	64	yes	-	-	-	-	-
▀ dct_Pipeline_WR_Loop_Row_WR_Loop_Col	-	-	-	-	67	670.000	-	67	-	no	0	0	32	156	0
WR_Loop_Row_WR_Loop_Col	-	-	-	-	65	650.000	3	1	64	yes	-	-	-	-	-
▀ Row_DCT_Loop	-	-	-	-	184	1.840E3	23	-	8	no	-	-	-	-	-
▀ dct_1d	-	-	-	-	21	210.000	-	21	-	no	0	8	658	345	0
▀ dct_1d_Pipeline_DCT_Outer_Loop	-	-	-	-	16	160.000	-	16	-	no	0	8	517	216	0
DCT_Outer_Loop	-	-	-	-	14	140.000	8	1	8	yes	-	-	-	-	-
Col_DCT_Loop	-	-	-	-	184	1.840E3	23	-	8	no	-	-	-	-	-

DCT_Outer_Loop flattening


Improve Memory Bandwidth

Create a new solution by copying the previous solution (Solution3) settings. Apply **ARRAY_PARTITION** directive to **buf_2d_in** of **dct** (since the bottleneck was on **src** port of the **dct_1d** function, which was passed via **in_block** of the **dct_2d** function, which in turn was passed via **buf_2d_in** of the **dct** function) and **col_inbuf** of **dct_2d**. Generate the solution.

1. Select **Project > New Solution** to create a new solution.
2. A *Solution Configuration* dialog box will appear. Click the **Finish** button (with Solution4 selected).
3. With *dct.c* open, select **buf_2d_in** array of the **dct** function in the Directive pane, right-click on it and select **Insert Directive...**

The **buf_2d_in** array is selected since the bottleneck was on **src** port of the **dct_1d** function, which was passed via **in_block** of the **dct_2d** function, which in turn was passed via **buf_2d_in** of the **dct** function).

4. A pop-up menu shows up listing various directives. Select **ARRAY_PARTITION** directive.
5. Make sure that the type is *complete*. Enter **2** in the dimension field and click **OK**.



*Applying **ARRAY_PARTITION** directive to memory buffer*

6. Similarly, apply the **ARRAY_PARTITION** directive with dimension of 2 to the **col_inbuf** array.
7. Click on the **Synthesis** button.
8. When the synthesis is completed, select **Project > Compare Reports...** to compare the two solutions.
9. Select **Solution4** and **Solution5** from the Available Reports, and click on the **Add>>** button.
10. Observe that the latency reduced from 643 to 579 clock cycles.

Performance Estimates			
Timing			
Clock		solution4	solution5
ap_clk	Target	10.00 ns	10.00 ns
	Estimated	6.508 ns	6.508 ns
Latency			
		solution4	solution5
Latency (cycles)	min	643	579
	max	643	579
Latency (absolute)	min	6.430 us	5.790 us
	max	6.430 us	5.790 us
Interval (cycles)	min	644	580
	max	644	580

Performance comparison after array partitioning

11. Scroll down in the comparison report to view the resources utilization. Observe the increase in the FF resource utilization (almost double).

Utilization Estimates		
	solution4	solution5
BRAM_18K	5	3
DSP	8	8
FF	827	1463
LUT	1377	1842
URAM	0	0

Resources utilization after array partitioning

12. Expand the Loop entry in the **Performance & Resource Estimates** in *Synthesis Summary* and observe that the Pipeline II is now 1.

Apply DATAFLOW Directive

Create a new solution by copying the previous solution (Solution5) settings. Apply the DATAFLOW directive to improve the throughput. Generate the solution and analyze the output.

1. Select **Project > New Solution**.
2. A *Solution Configuration* dialog box will appear. Click the **Finish** button (with Solution5 selected).
3. Close all inactive solution windows by selecting **Project > Close Inactive Solution Tabs**.
4. Select function **dct** in the directives pane, right-click on it and select **Insert Directive...**
5. Select **DATAFLOW** directive to improve the throughput.
6. Click on the *Synthesis* button.
7. When the synthesis is completed, the synthesis report is automatically opened.
8. Observe that dataflow type pipeline throughput is listed in the "Performance Estimates"

Performance Estimates

- Timing**
 - Summary**

Clock	Target	Estimated	Uncertainty
ap_clk	10.00 ns	6.508 ns	2.70 ns
- Latency**
 - Summary**

Latency (cycles)		Latency (absolute)		Interval (cycles)		Type
min	max	min	max	min	max	
577	577	5.770 us	5.770 us	443	443	dataflow
 - Detail**
 - Instance**
 - Loop**

Performance estimate after DATAFLOW directive applied

- The Dataflow pipeline throughput indicates the number of clock cycles between each set of inputs reads (interval parameter). If this value is less than the design latency it indicates the design can start processing new inputs before the currents input data are output.
- Note that the dataflow is only supported for the functions and loops at the top-level, not those which are down through the design hierarchy. Only loops and functions exposed at the toplevel of the design will get benefit from dataflow optimization.

9. Scrolling down into the *Utilization Estimates*, observe that the number of *BRAM_18K* required at the top-level remained at 3.

Utilization Estimates

- Summary**

Name	BRAM_18K	DSP	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	38	-
FIFO	-	-	-	-	-
Instance	2	8	1192	1486	0
Memory	1	-	256	32	0
Multiplexer	-	-	-	72	-
Register	-	-	8	-	-
Total	3	8	1456	1628	0
Available	280	220	106400	53200	0
Utilization (%)	1	3	1	3	0

Resource estimate with DATAFLOW directive applied

10. Look at the console view and notice that **dct_coeff_table** is automatically partitioned in dimension 2.

11. The `buf_2d_in` and `col_inbuf` arrays are partitioned as we had applied the directive in the previous run.
The dataflow is applied at the top-level which created channels between top-level functions `read_data`, `dct_2d`, and `write_data`.

```
INFO: [XFORM 203-712] Applying dataflow to function 'dct' (dct.c:81:1), detected/extracted 3 process function(s):
  'read_data.1'
  'dct_2d'
  'write_data.1'.
INFO: [XFORM 203-11] Balancing expressions in function 'dct_Id.4.5' (dct.c:6:24)...8 expression(s) balanced.
INFO: [HLS 200-111] Finished Loop, function and other optimizations: CPU user time: 0.08 seconds. CPU system time:
INFO: [XFORM 203-541] Flattening a loop nest 'WR_Loop_Row' (dct.c:68:8) in function 'write_data.1'.
INFO: [XFORM 203-541] Flattening a loop nest 'RD_Loop_Row' (dct.c:56:8) in function 'read_data.1'.
INFO: [XFORM 203-541] Flattening a loop nest 'Xpose_Row_Outer_Loop' (dct.c:28:16) in function 'dct_2d'.
INFO: [XFORM 203-541] Flattening a loop nest 'Xpose_Col_Outer_Loop' (dct.c:28:16) in function 'dct_2d'.
INFO: [HLS 200-472] Inferring partial write operation for 'buf_0' (dct.c:62:20)
INFO: [HLS 200-472] Inferring partial write operation for 'dst' (dct.c:19:14)
INFO: [HLS 200-472] Inferring partial write operation for 'col_inbuf' (dct.c:40:26)
INFO: [HLS 200-472] Inferring partial write operation for 'out_block' (dct.c:51:26)
INFO: [HLS 200-111] Finished Architecture Synthesis: CPU user time: 0.06 seconds. CPU system time: 0.01 seconds. El
INFO: [HLS 200-10] Starting hardware synthesis ...
INFO: [HLS 200-10] Synthesizing 'dct' ...
```

Console view of synthesis process after DATAFLOW directive applied

Perform performance analysis by switching to the Synthesis Summary and looking at the **dct Performance & Resource Estimates** view.

1. Switch to the *Synthesis Summary* perspective, expand the *Performance & Resource Estimates* entries, and select the **dct_2d** entry.
2. Observe that most of the latency and interval (throughput) is caused by the `dct_2d` function. The interval of the top-level function `dct`, is less than the sum of the intervals of the `read_data`, `dct_2d`, and `write_data` functions indicating that they operate in parallel and `dct_2d` is the limiting factor. It can be seen that `dct_2d` is not completely operating in parallel as `Row_DCT_Loop` and `Col_DCT_Loop` were not pipelined.

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
↳ dct			.	.	577	5.770E3	.	443	.	dataflow	3	8	1456	1628	0
↳ read_data_1			.	66	660.000	.	66	.	no	0	0	25	146	0	
↳ RD_Loop_Row_RD_Loop_Col			.	64	640.000	.	64	.	yes	-	-	-	-	-	
↳ dct_2d			.	442	4.420E3	.	442	.	no	2	8	1131	1180	0	
↳ dct_2d_Pipeline_Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop			.	67	670.000	.	67	.	no	0	0	34	156	0	
↳ dct_2d_Pipeline_Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop			.	67	670.000	.	67	.	no	0	0	38	170	0	
↳ Row_DCT_Loop			.	152	1.520E3	19	-	8	no	-	-	-	-	-	
↳ dct_id_4_5			.	16	160.000	.	16	.	no	0	8	520	216	0	
↳ Col_DCT_Loop			.	152	1.520E3	19	-	8	no	-	-	-	-	-	
↳ write_data_1			.	67	670.000	.	67	.	no	0	0	36	160	0	
↳ WR_Loop_Row_WR_Loop_Col			.	65	650.000	3	1	64	yes	-	-	-	-	-	

Performance analysis after the DATAFLOW directive

One of the limitations of the dataflow optimization is that it only works on top-level loops and functions. One way to have the blocks in `dct_2d` operate in parallel would be to pipeline the entire function. This however would unroll all the loops and can sometimes lead to a large area increase. An alternative is to raise these loops up to the top-level of hierarchy, where dataflow optimization can be applied, by removing the `dct_2d` hierarchy, i.e. inline the `dct_2d` function.

Apply INLINE Directive

Create a new solution by copying the previous solution (Solution6) settings. Apply INLINE directive to `dct_2d`. Generate the solution and analyze the output.

1. Select **Project > New Solution**.
2. A *Solution Configuration* dialog box will appear. Click the **Finish** button (with *Solution6* selected).
3. Select the function *dct_2d* in the directives pane, right-click on it and select **Insert Directive...**
4. A pop-up menu shows up listing various directives. Select **INLINE** directive. The INLINE directive causes the function to which it is applied to be inlined: its hierarchy is dissolved.
5. Click on the *Synthesis* button.
6. When the synthesis is completed, the synthesis report will be opened.
7. Observe that the latency reduced from 577 to 416 clock cycles, and the Dataflow pipeline throughput drastically reduced from 443 to 73 clock cycles.
8. Examine the synthesis log to see what transformations were applied automatically.
 - The *dct_1d* function calls are now automatically inlined into the loops from which they are called, which allows the loop nesting to be flattened automatically.
 - Note also that the DSP48E usage has doubled (from 8 to 16). This is because, previously a single instance of *dct_1d* was used to do both row and column processing; now that the row and column loops are executing concurrently, this can no longer be the case and two copies of *dct_1d* are required: Vitis HLS will seek to minimize the number of clocks, even if it means increasing the area.

```

INFO: [XFORM 203-712] Applying dataflow to function 'dct' (dct.c:26:1), detected/extracted 6 process function(s):
  'read_data.1'
  'Loop_Row_DCT_Loop_proc'
  'Loop_Xpose_Row_Outer_Loop_proc'
  'Loop_Col_DCT_Loop_proc'
  'Loop_Xpose_Col_Outer_Loop_proc'
  'write_data.1'.
INFO: [XFORM 203-602] Inlining function 'dct_1d' into 'Loop_Row_DCT_Loop_proc' (dct.c:33) automatically.
INFO: [XFORM 203-602] Inlining function 'dct_1d' into 'Loop_Col_DCT_Loop_proc' (dct.c:44) automatically.
INFO: [XFORM 203-11] Balancing expressions in function 'Loop_Row_DCT_Loop_proc' (dct.c:6:24)...8 expression(s) balanced.
INFO: [XFORM 203-11] Balancing expressions in function 'Loop_Col_DCT_Loop_proc' (dct.c:6:24)...8 expression(s) balanced.

```

Console view after INLINE directive applied to dct_2d

9. Switch to the *Synthesis Summary* perspective, expand the *Performance & Resource Estimates* entries, and select the **dct** entry.

Observe that the *dct_2d* entry is now replaced with *dct_Loop_Row_DCT_Loop_proc*, *dct_Loop_Xpose_Row_Outer_Loop_proc*, *dct_Loop_Col_DCT_Loop_proc*, and *dct_Loop_Xpose_Col_Outer_Loop_proc* since the *dct_2d* function is inlined. Also observe that all the functions are operating in parallel, yielding the top-level function interval (throughput) of 73 clock cycles.

Performance & Resource Estimates ⓘ															
Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LLUT	URAM
↳ dct	-	-	416	4.160E3	-	-	73	-	-	dataflow	3	16	1271	1532	0
↳ read_data_1	-	-	66	660.000	-	66	-	-	no	0	0	25	146	0	
↳ RD_Loop_Row_RD_Loop_Col	-	-	64	640.000	2	1	64	yes	-	-	-	-	-	-	
↳ Loop_Row_DCT_Loop_proc	-	-	72	720.000	-	72	-	-	no	0	8	305	305	0	
↳ Row_DCT_Loop_DCT_Outer_Loop	-	-	70	700.000	8	1	64	yes	-	-	-	-	-	-	
↳ Loop_Xpose_Row_Outer_Loop_proc	-	-	67	670.000	-	67	-	-	no	0	0	34	160	0	
↳ Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	-	65	650.000	3	1	64	yes	-	-	-	-	-	-	
↳ Loop_Col_DCT_Loop_proc	-	-	72	720.000	-	72	-	-	no	0	8	305	305	0	
↳ Col_DCT_Loop_DCT_Outer_Loop	-	-	70	700.000	8	1	64	yes	-	-	-	-	-	-	
↳ Loop_Xpose_Col_Outer_Loop_proc	-	-	67	670.000	-	67	-	-	no	0	0	38	174	0	
↳ Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	-	65	650.000	3	1	64	yes	-	-	-	-	-	-	
↳ write_data_1	-	-	67	670.000	-	67	-	-	no	0	0	36	160	0	
↳ WR_Loop_Row_WR_Loop_Col	-	-	65	650.000	3	1	64	yes	-	-	-	-	-	-	

Performance analysis after the **INLINE** directive

10. Close Vitis HLS by selecting **File > Exit**.

Conclusion

In this lab, you learned various techniques to improve the performance and balance resource utilization. PIPELINE directive when applied to outer loop will automatically cause the inner loop to unroll. When a loop is unrolled, resources utilization increases as operations are done concurrently. Partitioning memory may improve performance but will increase BRAM utilization. When INLINE directive is applied to a function, the lower level hierarchy is automatically dissolved. When DATAFLOW directive is applied, the default memory buffers (of ping-pong type) are automatically inserted between the top-level functions and loops. The console logs can provide insight on what is going on.

Answers

Answers for question 1:

Estimated clock period: **6.508 ns**

Worst case latency: **423 clock cycles**

Number of DSP48E used: **17**

Number of BRAMs used: **16**

Number of FFs used: **929**

Number of LUTs used: **2012**

Copyright© 2022, Advanced Micro Devices, Inc.